Assignment 2

Kevin Gardner

2/15/2022

Following is the link to my GitHub account:

https://github.com/Kgardner22/64060_-kgardner

Import and Prepare Data:

Import the UniversalBank.csv file

```
UniversalBank <- read.table('C:/R/MyData/UniversalBank.csv', header = T, sep = ',')

summary(UniversalBank)
```

```
##        ID            Age          Experience       Income                ZIP.Code
##  Min.   :   1   Min.   :23.00   Min.   :-3.0   Min.   :  8.00   Min.   :
9307
##  1st Qu.:1251   1st Qu.:35.00   1st Qu.:10.0   1st Qu.: 39.00   1st
Qu.:91911
##  Median :2500   Median :45.00   Median :20.0   Median : 64.00   Median
:93437
##  Mean   :2500   Mean   :45.34   Mean   :20.1   Mean   : 73.77   Mean
:93153
##  3rd Qu.:3750   3rd Qu.:55.00   3rd Qu.:30.0   3rd Qu.: 98.00   3rd
Qu.:94608
##  Max.   :5000   Max.   :67.00   Max.   :43.0   Max.   :224.00   Max.
:96651
##      Family          CCAvg          Education        Mortgage
##  Min.   :1.000   Min.   : 0.000   Min.   :1.000   Min.   :  0.0
##  1st Qu.:1.000   1st Qu.: 0.700   1st Qu.:1.000   1st Qu.:  0.0
##  Median :2.000   Median : 1.500   Median :2.000   Median :  0.0
##  Mean   :2.396   Mean   : 1.938   Mean   :1.881   Mean   : 56.5
##  3rd Qu.:3.000   3rd Qu.: 2.500   3rd Qu.:3.000   3rd Qu.:101.0
```

```
##  Max.   :4.000   Max.   :10.000   Max.   :3.000   Max.   :635.0
##  Personal.Loan   Securities.Account   CD.Account         Online
##  Min.   :0.000   Min.   :0.0000     Min.   :0.0000   Min.   :0.0000
##  1st Qu.:0.000   1st Qu.:0.0000     1st Qu.:0.0000   1st Qu.:0.0000
##  Median :0.000   Median :0.0000     Median :0.0000   Median :1.0000
##  Mean   :0.096   Mean   :0.1044     Mean   :0.0604   Mean   :0.5968
##  3rd Qu.:0.000   3rd Qu.:0.0000     3rd Qu.:0.0000   3rd Qu.:1.0000
##  Max.   :1.000   Max.   :1.0000     Max.   :1.0000   Max.   :1.0000
##    CreditCard
##  Min.   :0.000
##  1st Qu.:0.000
##  Median :0.000
##  Mean   :0.294
##  3rd Qu.:1.000
##  Max.   :1.000
```

Create a copy of the original data file to preserve

```
Original_File <- UniversalBank
```

REQUIREMENT 1:

Transform categorical predictors with more than two categories into dummy variables FIRST. Need to do this for 'Education' and 'Personal.Loan'.

```
library(caret)

## Loading required package: ggplot2

## Loading required package: lattice

library(class)

# Remove unnecessary attributes
UniversalBank$ID <- NULL
UniversalBank$ZIP.Code <- NULL

# Transform to factors
UniversalBank$Education=as.factor(UniversalBank$Education)
UniversalBank$Personal.Loan=as.factor(UniversalBank$Personal.Loan)

# Use dummyVars function to create a model
dummies <- dummyVars(Personal.Loan ~ ., data = UniversalBank)
UniversalBank_dummy <- as.data.frame(predict(dummies, newdata =
UniversalBank))

## Warning in model.frame.default(Terms, newdata, na.action = na.action, xlev
=
## object$lvls): variable 'Personal.Loan' is not a factor
```

Normalize the data (I'm using z-score scaling as input method)

```
Norm_model <- preProcess(UniversalBank_dummy, method = c("center", "scale"))
UniversalBank_norm <- predict(Norm_model, UniversalBank_dummy)
summary(UniversalBank_norm)
```

```
##       Age              Experience           Income             Family       
##  Min.   :-1.94871   Min.   :-2.014710   Min.   :-1.4288   Min.   :-1.2167  
##  1st Qu.:-0.90188   1st Qu.:-0.881116   1st Qu.:-0.7554   1st Qu.:-1.2167  
##  Median :-0.02952   Median :-0.009121   Median :-0.2123   Median :-0.3454  
##  Mean   : 0.00000   Mean   : 0.000000   Mean   : 0.0000   Mean   : 0.0000  
##  3rd Qu.: 0.84284   3rd Qu.: 0.862874   3rd Qu.: 0.5263   3rd Qu.: 0.5259  
##  Max.   : 1.88967   Max.   : 1.996468   Max.   : 3.2634   Max.   : 1.3973  
##      CCAvg          Education.1       Education.2       Education.3      
##  Min.   :-1.1089   Min.   :-0.8495   Min.   :-0.6245   Min.   :-0.6549  
##  1st Qu.:-0.7083   1st Qu.:-0.8495   1st Qu.:-0.6245   1st Qu.:-0.6549  
##  Median :-0.2506   Median :-0.8495   Median :-0.6245   Median :-0.6549  
##  Mean   : 0.0000   Mean   : 0.0000   Mean   : 0.0000   Mean   : 0.0000  
##  3rd Qu.: 0.3216   3rd Qu.: 1.1770   3rd Qu.: 1.6010   3rd Qu.: 1.5266  
##  Max.   : 4.6131   Max.   : 1.1770   Max.   : 1.6010   Max.   : 1.5266  
##     Mortgage       Securities.Account   CD.Account          Online       
##  Min.   :-0.5555   Min.   :-0.3414    Min.   :-0.2535   Min.   :-1.2165  
##  1st Qu.:-0.5555   1st Qu.:-0.3414    1st Qu.:-0.2535   1st Qu.:-1.2165  
##  Median :-0.5555   Median :-0.3414    Median :-0.2535   Median : 0.8219  
##  Mean   : 0.0000   Mean   : 0.0000    Mean   : 0.0000   Mean   : 0.0000  
##  3rd Qu.: 0.4375   3rd Qu.:-0.3414    3rd Qu.:-0.2535   3rd Qu.: 0.8219  
##  Max.   : 5.6875   Max.   : 2.9286    Max.   : 3.9438   Max.   : 0.8219  
##    CreditCard      
##  Min.   :-0.6452  
##  1st Qu.:-0.6452  
##  Median :-0.6452  
##  Mean   : 0.0000  
##  3rd Qu.: 1.5495  
##  Max.   : 1.5495
```

Add back in the target variable (Personal.Loan)

```
UniversalBank_norm$Personal.Loan <- UniversalBank$Personal.Loan
```

We need to divide the data into training (60%) and validation (40%) sets

```
Train_Index <- createDataPartition(UniversalBank$Personal.Loan, p=0.6, list = FALSE)
Train.df <- UniversalBank_norm[Train_Index,]
Validation.df <- UniversalBank_norm[-Train_Index,]
```

Create data frame with values to predict Age = 40, Experience = 10, Income = 84, Family = 2, CCAvg = 2, Education_1 = 0, Education_2 = 1, Education_3 = 0, Mortgage = 0, Securities Account = 0, CD Account = 0, Online = 1, and Credit Card = 1.

```
To_Predict <- data.frame (Age=40, Experience=10, Income=84, Family=2, CCAvg=2, Education.1=0, Education.2=1, Education.3=0, Mortgage=0, Securities.Account=0, CD.Account=0, Online=1, CreditCard=1)
```

```
print(To_Predict)
```

```
##   Age Experience Income Family CCAvg Education.1 Education.2 Education.3
## 1  40         10     84      2     2           0           1           0
##   Mortgage Securities.Account CD.Account Online CreditCard
## 1        0                  0          0      1          1
```

Normalize this new record (To_Predict) using the same model we applied to the original dataset

```
To_Predict_norm <- predict(Norm_model, To_Predict)
print(To_Predict_norm)
```

```
##          Age Experience    Income     Family     CCAvg Education.1
Education.2
## 1 -0.4657003 -0.8811162 0.2221371 -0.3453975 0.0355115  -0.8494814
1.601024
##   Education.3   Mortgage Securities.Account CD.Account    Online
CreditCard
## 1  -0.6548999 -0.5554684         -0.3413892 -0.2535149 0.8218687
1.549477
```

Use k-NN function to make the prediction.

Perform a k-NN classification with all predictors EXCEPT ID and Zip_Code using k=1. Specify success class as 1 (loan acceptance) and use the default cutoff value of 0.5.

```
Prediction <- knn(train = Train.df[,1:13],
                  test = To_Predict_norm[,1:13],
                  cl=Train.df$Personal.Loan,
                  k=1)
print(Prediction)
```

```
## [1] 0
## Levels: 0 1
```

ANSWER - REQUIREMENT 1:

As shown in the above results (with k=1), the prediction for this observation is that Personal.Loan = 0 meaning, this individual is predicted to NOT accept the personal loan being offered.

REQUIREMENT 2:

What is a choice of k that balances between overfitting and ignoring the predictor information?

```
set.seed(123)

fitControl <- trainControl(method = "repeatedcv", number = 3, repeats = 2)
```

```
searchGrid <- expand.grid(k = 1:15)

Knn.model <- train(Personal.Loan ~ .,
                   data = Train.df,
                   method = 'knn',
                   tuneGrid = searchGrid,
                   trControl = fitControl,)

Knn.model
```

```
## k-Nearest Neighbors
##
## 3000 samples
##   13 predictor
##    2 classes: '0', '1'
##
## No pre-processing
## Resampling: Cross-Validated (3 fold, repeated 2 times)
## Summary of sample sizes: 2000, 2000, 2000, 2000, 2000, 2000, ...
## Resampling results across tuning parameters:
##
##   k   Accuracy   Kappa
##    1  0.9605000  0.7478112
##    2  0.9548333  0.7089392
##    3  0.9576667  0.7112729
##    4  0.9538333  0.6774099
##    5  0.9556667  0.6864548
##    6  0.9546667  0.6786735
##    7  0.9516667  0.6488616
##    8  0.9505000  0.6400390
##    9  0.9496667  0.6282323
##   10  0.9471667  0.6091686
##   11  0.9461667  0.5956549
##   12  0.9453333  0.5877380
##   13  0.9446667  0.5810143
##   14  0.9438333  0.5732348
##   15  0.9436667  0.5706980
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was k = 1.
```

ANSWER - REQUIREMENT 2:

The above Knn.model indicates that the k value with the highest accuracy is a value of k=3

REQUIREMENT 3:

Show the confusion matrix for the validation data that results from using the best k.

Use the predict function of the caret package to make predictions on the validation set.

```
predictions <- predict(Knn.model, Validation.df)
```

Compare predictions from the Knn.model to the actual Personal.Loan labels in the validation set to compute the confusion matrix

```
confusionMatrix(predictions, Validation.df$Personal.Loan)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##          0 1783   69
##          1   25  123
##
##                Accuracy : 0.953
##                  95% CI : (0.9428, 0.9619)
##     No Information Rate : 0.904
##     P-Value [Acc > NIR] : 2.260e-16
##
##                   Kappa : 0.6983
##
##  Mcnemar's Test P-Value : 9.202e-06
##
##             Sensitivity : 0.9862
##             Specificity : 0.6406
##          Pos Pred Value : 0.9627
##          Neg Pred Value : 0.8311
##              Prevalence : 0.9040
##          Detection Rate : 0.8915
##    Detection Prevalence : 0.9260
##       Balanced Accuracy : 0.8134
##
##        'Positive' Class : 0
##
```

REQUIREMENT 4:

Consider the following customer: Age = 40, Experience = 10, Income = 84, Family = 2, CCAvg = 2, Education_1 = 0, Education_2 = 1, Education_3 = 0, Mortgage = 0, Securities Account = 0, CD Account = 0, Online = 1 and Credit Card = 1. Classify the customer using the best k.

```
# The independent variables specified in this requirement are the same as
those specified in requirement 1. Therefore, this data.frame is already build
and normalized (To_Predict_norm).

print(To_Predict)
```

```
##   Age Experience Income Family CCAvg Education.1 Education.2 Education.3
## 1  40         10     84      2     2           0           1           0
```

```
##   Mortgage Securities.Account CD.Account Online CreditCard
## 1        0                  0          0      1          1
```

Using the normalized prediction file (To_Predict_norm), we will use the Knn.model to predict using the best k value (k=3)

```
predict(Knn.model, To_Predict_norm)
```

```
## [1] 0
## Levels: 0 1
```

ANSWER - REQUIREMENT 4:

Using the best k value, the above results (with k=3) is predicting this observation will have Personal.Loan = 0 meaning, this individual is predicted to NOT accept the personal loan being offered.